

```

from models import resnet15v2
from train import get_dataloaders, evaluate
import torch
import torch.nn as nn
from torchsummary import summary

print(torch.cuda.is_available())
print(torch.cuda.device_count())

True
1

device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")
print(device)

cuda:0

```

Adam Optimizer with 50 Epochs

```

# Load the model
model = resnet15v2()
model.load_state_dict(torch.load('expt_50_adam/best.pth',
map_location=device))
model = model.to(device)

```

Model Architecture

```
summary(model, (2, 32, 32))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	1,152
BatchNorm2d-2	[-1, 64, 32, 32]	128
ReLU-3	[-1, 64, 32, 32]	0
Conv2d-4	[-1, 64, 32, 32]	36,864
BatchNorm2d-5	[-1, 64, 32, 32]	128
ReLU-6	[-1, 64, 32, 32]	0
Identity-7	[-1, 64, 32, 32]	0
Conv2d-8	[-1, 64, 32, 32]	36,864
BatchNorm2d-9	[-1, 64, 32, 32]	128
ReLU-10	[-1, 64, 32, 32]	0
Conv2d-11	[-1, 64, 32, 32]	36,864
BatchNorm2d-12	[-1, 64, 32, 32]	128
ReLU-13	[-1, 64, 32, 32]	0
BasicBlock-14	[-1, 64, 32, 32]	0
Conv2d-15	[-1, 64, 32, 32]	36,864
BatchNorm2d-16	[-1, 64, 32, 32]	128

ReLU-17	[-1, 64, 32, 32]	0
Conv2d-18	[-1, 64, 32, 32]	36,864
BatchNorm2d-19	[-1, 64, 32, 32]	128
ReLU-20	[-1, 64, 32, 32]	0
BasicBlock-21	[-1, 64, 32, 32]	0
Conv2d-22	[-1, 128, 16, 16]	73,728
BatchNorm2d-23	[-1, 128, 16, 16]	256
ReLU-24	[-1, 128, 16, 16]	0
Conv2d-25	[-1, 128, 16, 16]	147,456
BatchNorm2d-26	[-1, 128, 16, 16]	256
Conv2d-27	[-1, 128, 16, 16]	8,192
BatchNorm2d-28	[-1, 128, 16, 16]	256
ReLU-29	[-1, 128, 16, 16]	0
BasicBlock-30	[-1, 128, 16, 16]	0
Conv2d-31	[-1, 128, 16, 16]	147,456
BatchNorm2d-32	[-1, 128, 16, 16]	256
ReLU-33	[-1, 128, 16, 16]	0
Conv2d-34	[-1, 128, 16, 16]	147,456
BatchNorm2d-35	[-1, 128, 16, 16]	256
ReLU-36	[-1, 128, 16, 16]	0
BasicBlock-37	[-1, 128, 16, 16]	0
Conv2d-38	[-1, 256, 8, 8]	294,912
BatchNorm2d-39	[-1, 256, 8, 8]	512
ReLU-40	[-1, 256, 8, 8]	0
Conv2d-41	[-1, 256, 8, 8]	589,824
BatchNorm2d-42	[-1, 256, 8, 8]	512
Conv2d-43	[-1, 256, 8, 8]	32,768
BatchNorm2d-44	[-1, 256, 8, 8]	512
ReLU-45	[-1, 256, 8, 8]	0
BasicBlock-46	[-1, 256, 8, 8]	0
Conv2d-47	[-1, 256, 8, 8]	589,824
BatchNorm2d-48	[-1, 256, 8, 8]	512
ReLU-49	[-1, 256, 8, 8]	0
Conv2d-50	[-1, 256, 8, 8]	589,824
BatchNorm2d-51	[-1, 256, 8, 8]	512
ReLU-52	[-1, 256, 8, 8]	0
BasicBlock-53	[-1, 256, 8, 8]	0
Identity-54	[-1, 256, 8, 8]	0
AdaptiveAvgPool2d-55	[-1, 256, 1, 1]	0
Linear-56	[-1, 2]	514

```

=====
Total params: 2,812,034
Trainable params: 2,812,034
Non-trainable params: 0

```

```

-----
Input size (MB): 0.01
Forward/backward pass size (MB): 16.63
Params size (MB): 10.73

```

Estimated Total Size (MB): 27.36

```
-----  
train_loader, test_loader = get_dataloaders(batch_size=512)  
criterion = nn.BCEWithLogitsLoss()
```

Evaluation

```
test_acc, test_loss = evaluate(model, criterion, test_loader, device)  
print('Test accuracy: ', test_acc*100, '%')  
print('Test loss: ', test_loss)
```

Test accuracy: 50.43473895582329 %
Test loss: 1.015372067187206

Adam Optimizer with 20 Epochs

```
model2 = resnet15v2()  
model2.load_state_dict(torch.load('expt_20_adam/best.pth',  
map_location=device))  
model2 = model2.to(device)
```

Evaluation

```
test_acc2, test_loss2 = evaluate(model2, criterion, test_loader,  
device)  
print('Test accuracy: ', test_acc2*100, '%')  
print('Test loss: ', test_loss2)
```

Test accuracy: 75.47791164658635 %
Test loss: 0.5066263087111783

AdamW optimizer with 20 Epochs

```
model3 = resnet15v2()  
model3.load_state_dict(torch.load('expt_20_adamw/best.pth',  
map_location=device))  
model3 = model3.to(device)
```

Evaluation

```
test_acc3, test_loss3 = evaluate(model3, criterion, test_loader,  
device)  
print('Test accuracy: ', test_acc3*100, '%')  
print('Test loss: ', test_loss3)
```

Test accuracy: 74.68072289156626 %
Test loss: 0.5188913260885032

Plot

```
import matplotlib.pyplot as plt

test_accs = [test_acc * 100, test_acc2 * 100, test_acc3 * 100]
test_losses = [test_loss, test_loss2, test_loss3]
model_names = ['expt_50_adam', 'expt_20_adam', 'expt_20_adamw']

fig, ax1 = plt.subplots(figsize=(8,6))

color_acc = 'tab:blue'
ax1.set_xlabel('Model')
ax1.set_ylabel('Test Accuracy (%)', color=color_acc)
bars = ax1.bar(model_names, test_accs, color=color_acc, alpha=0.6)
ax1.tick_params(axis='y', labelcolor=color_acc)
ax1.set_ylim(0, 100)

for bar in bars:
    height = bar.get_height()
    ax1.annotate(f'{height:.1f}%',
                xy=(bar.get_x() + bar.get_width() / 2, height),
                xytext=(0, 3),
                textcoords="offset points",
                ha='center', va='bottom')

ax2 = ax1.twinx()
color_loss = 'tab:red'
ax2.set_ylabel('Test Loss', color=color_loss)
ax2.plot(model_names, test_losses, color=color_loss, marker='o',
        linestyle='--', linewidth=2)
ax2.tick_params(axis='y', labelcolor=color_loss)

plt.title('Comparison of Test Accuracy and Loss for Three Models')
plt.show()
```

